

# COMP 8005 Assignment 01

---

Measuring the performance and efficiency of using multiple processes and threads.

**Thilina Ratnayake – A00802338**

**1/26/2015**

An experiment to observe the difference in performance between multi-threading versus multi-processing.

**Table Of Contents:**

<u>Page</u>	<u>Item</u>
2	Objective
2	<b>REPORT &amp; EXPERIMENT DESIGN</b>
2	Controls
3	Results
6	Observations and Remarks
6	Conclusion
	<b>PROGRAM DESIGN</b>
7	Finite State Machines
8	Program Components
9	Program Structure
10	Psuedo-Code: Processes
11	Psuedo-Code: Threads
12	Execution Instructions
13	<b>TESTING</b>
14	<b>Appendix A: Results in tabular form.</b>

**Objective:**

Use multiple processes and threads on the Linux operating system and measure the differences of performance (if any) between each mechanism.

**Experiment Design**

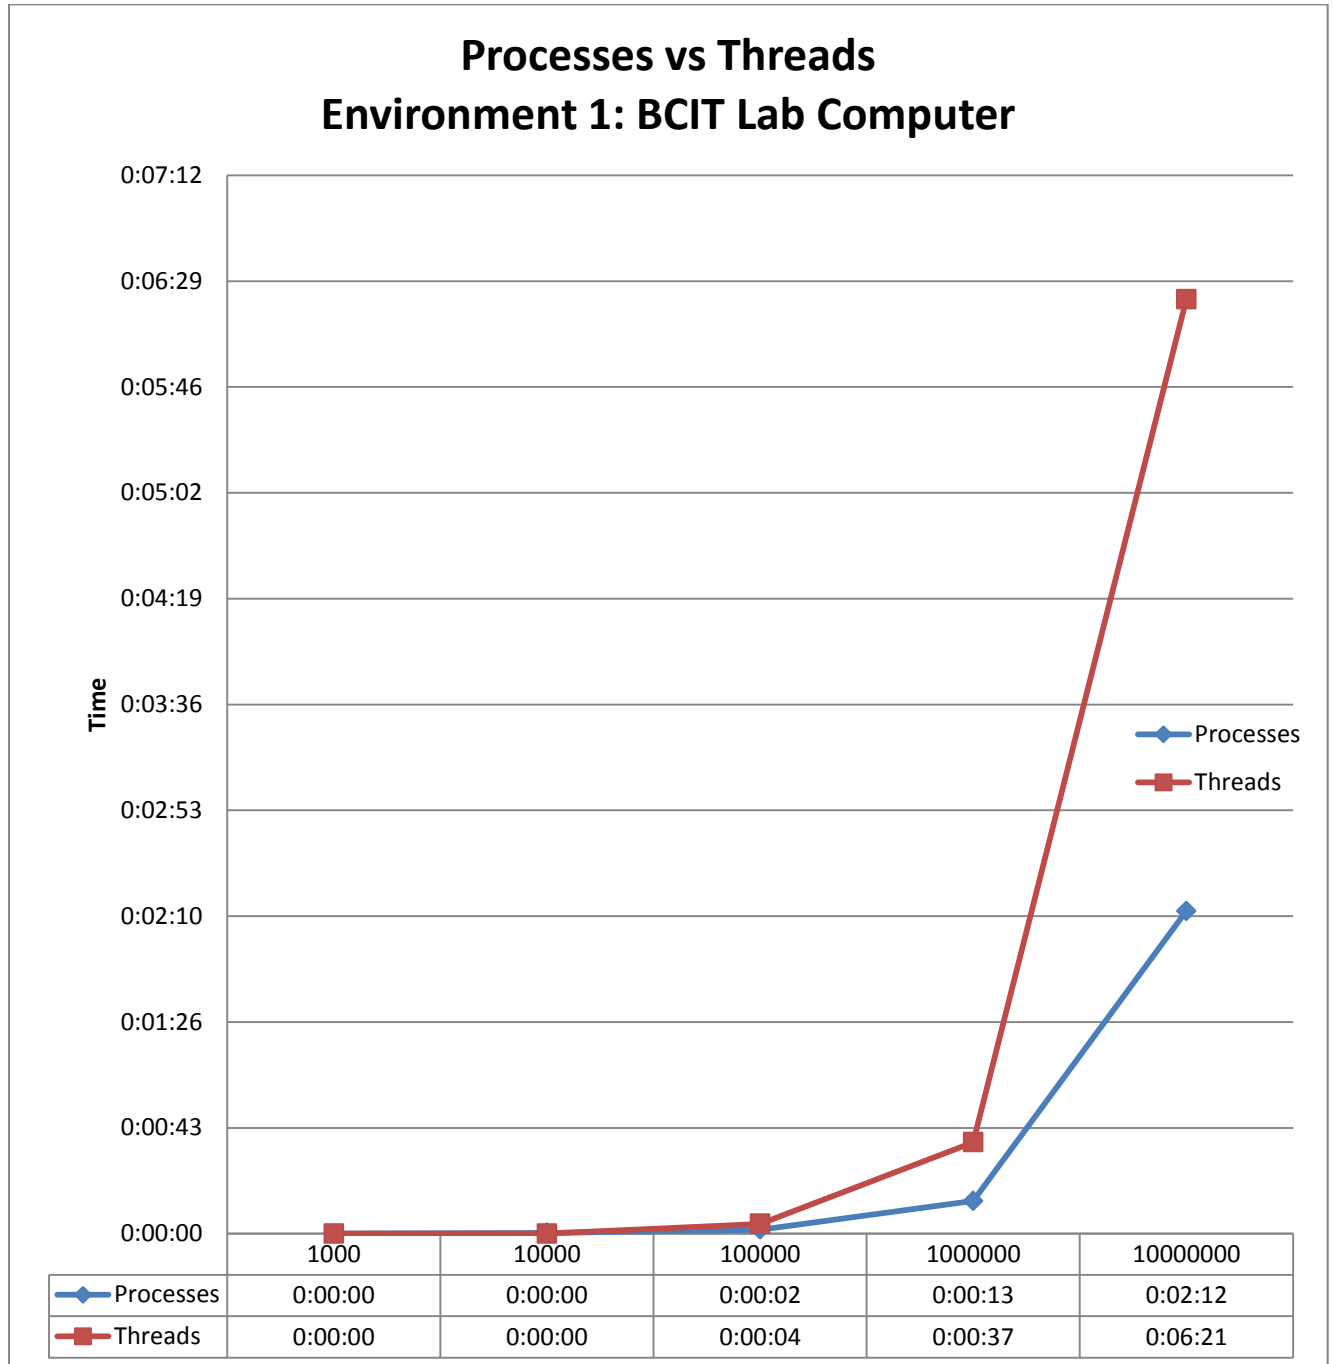
1. The experiment will measure the performance of each subject (processes or threads) based on time taken to finish a task.
2. The language that will be used is C, and will be written for and tested on the Linux OS.
3. The experiment involves both Mathematical Computations and file I/O.
  - a. Mathematical computations: A string ("HASHSTRINGX" where X will be the iteration number) will be hashed using the SHA1 hashing algorithm a specified number of times.
  - b. File I/O: The plaintext and hashed text will be printed onto text files.
4. Both subjects will be tested on hashing 1000,10000,100000,1000000, and 10,000,000 times.


**Controls:**

1. The experiment will be run 3 times for each subject, and completion times will be averaged.
2. The experiment will be run on the Fedora 21 distro.
3. The experiment will be run on two sets of hardware/environments.


**Specifications:**

<b>Environment 1 - BCIT Lab</b>	<b>Environment 2 -Personal PC with Linux VM</b>
<ul style="list-style-type: none"><li>• Processor Cores: 4</li><li>• Processor: Intel Core i5-2400k CPU @ 3.10GHz -64bit</li><li>• RAM:8GB</li></ul>	<ul style="list-style-type: none"><li>• Processor Cores: 4</li><li>• Processor: Intel Core i5-4670K CPU @ 3.40 GHz -64bit</li><li>• Processor Speed:</li><li>• RAM:1GB</li></ul>

**Results:****BCIT Lab:****Fig 1A: Chart of performance times over amount of iterations hashed in Environment 1(BCIT Lab)**

Linux VM:

**Fig 1C:** Chart of performance times over amount of iterations hashed in Environment 2 (Linux VM)

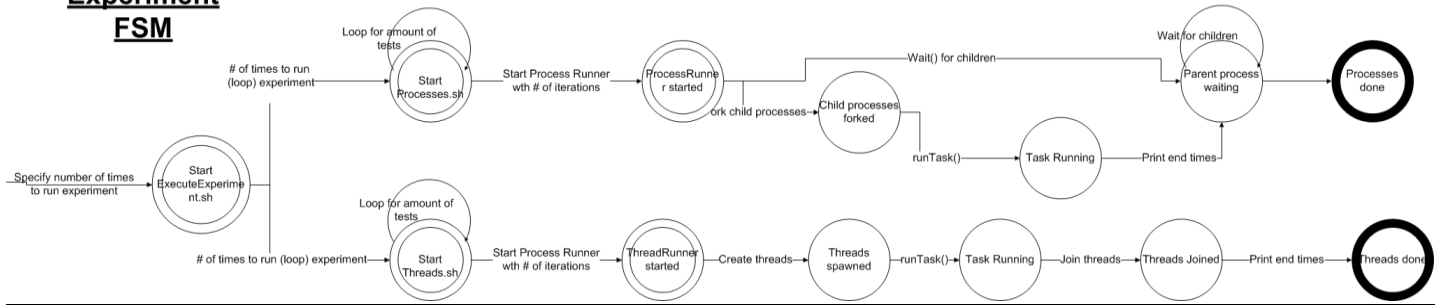

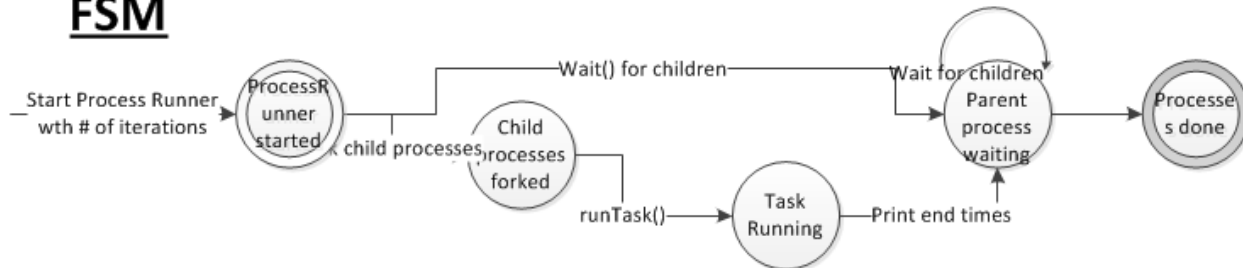
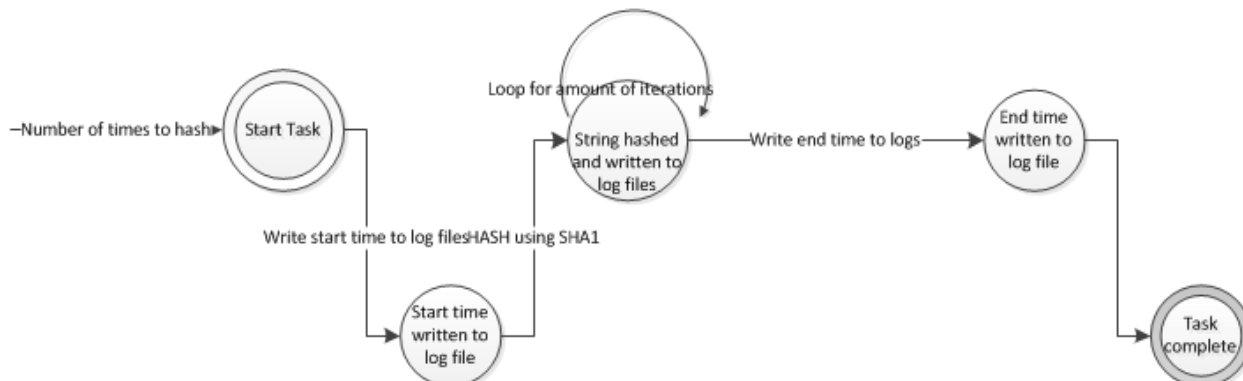
Final Comparison**Fig 1E:** Chart comparing results from both environments.

**Observations & Remarks:**

1. The performance between the two mechanisms are the same when dealing with a task of requiring lower processing requirements, as demonstrated by their performance from iterations 1000 to 100,000.
2. As compared to my colleague's experiments which utilized prime number decomposition as it's task instead of hashing, their results indicated that the difference in performance between the two mechanisms was marginal through-out the experiment, whereas in this experiment there is a drastic change in performance.
3. In trying to understand why multi-threading took so much time, the sched\_yield() call and mutex\_locks were taken out as they do not utilize any of the same resources, thus avoiding any race conditions. This change in the program only reduced ~30 seconds off the total time.

**Conclusion**

The conclusion that can be reached from this experiment is that the efficiency of utilizing a multi-threading mechanism decreases with the rise of processing requirements for a specific task. When compared to a multi-processing mechanism, which is approximately 291% more efficient than multi-threading at a task of 10,000,000 iterations.

**DESIGN****Finite State Diagrams***(These diagrams can be found in /FSMs for closer inspection)***Main  
Experiment  
FSM****Threads  
FSM****Processes  
FSM****Task  
FSM**



**Program Components:**

Component	Functions	Inputs	Outputs
ExecuteExperiment.sh	Script that runs Processes.sh and Threads.sh one after the other.	Number of times to run tests	--
Processes.sh	Script that runs ProcessRunner.c in a loop with desired amount of hash iterations.	Number of times to run tests (ProcessRunner)	./logfiles/DATETIME/DATETIME.csv
Threads.sh	Script that runs ThreadRunner.c in a loop with desired amount of hash iterations.	Number of times to run tests (ThreadRunner)	./logfiles/DATETIME/DATETIME.csv
ProcessRunner.c	Spawns 5 child processes and executes a task that hashes a string for the specified amount of iterations and outputs to logfiles.	Number of iterations.	./ProcessFiles/ProcessTaskOutputFileX.txt (where X = Process number)  <i>Note:</i> When run with the script, all output files will be deleted as part of tear down procedure to save disk space.
ThreadRunner.c	Spawns 5 child threads and executes a task that hashes a string for the specified amount of iterations and	Number of iterations.	./ThreadFiles/ThreadTaskOutputFileThreadX.txt (where X = Thread Number)  <i>Note:</i> When run with the script, all output files will be deleted as part of tear down procedure to save disk space.

**Program Structure:**

```
bash: $. command not found...
[root@localhost 8005-Assignment-1]# tree
.
├── ExecuteExperiment.sh
├── Processes
│   ├── logfiles
│   │   ├── 2015_01_23_22_29_59
│   │   │   ├── Processes2015_01_23_22_29_59.csv
│   │   │   └── Processes2015_01_23_22_32_52.csv
│   │   └── 2015_01_24_10_05_38
│   │       ├── Processes2015_01_24_10_05_38.csv
│   │       └── Processes2015_01_24_10_08_32.csv
│   ├── Makefile
│   ├── Processes
│   ├── Processes.sh
│   ├── ProcessFiles
│   ├── ProcessLogFile.txt
│   ├── ProcessRunner
│   └── ProcessRunner.c
├── PsuedoCode:\ Processes.txt
├── PsuedoCode:\ Tasks.txt
├── Threads
│   ├── logfiles
│   │   ├── 2015_01_23_22_35_46
│   │   │   ├── Threads2015_01_23_22_35_46.csv
│   │   │   └── Threads2015_01_23_22_43_23.csv
│   │   └── 2015_01_24_10_11_28
│   │       ├── Threads2015_01_24_10_11_28.csv
│   │       └── Threads2015_01_24_10_19_21.csv
│   ├── Makefile
│   ├── ThreadFiles
│   ├── ThreadLogFile.txt
│   ├── ThreadRunner
│   ├── ThreadRunner.c
│   └── Threads.sh
```

**Pseudo-Code****Processes:**

```
1  PsuedoCode: Processes
2
3  ▼ Main(number of processes, number of iterations){
4      - Write to main log file the starting time.
5
6  ▼      for loop (i to number of threads){
7          //fork processes.
8
9          pid = fork();
10     }
11
12     //handle the different type of processes
13
14  ▼     switch(pid){
15         if Unsuccessful forks:
16             error
17
18         if its a child process
19             case 0:
20                 run the (task)(number of iterations);
21
22         if its a parent
23     ▼     case(anything else):
24
25         loop through number of processes
26     ▼     {
27         wait for each process
28         //this next statement only executes if the child is done.
29         subtract -1 number of processes.
30     }
31
32     //All children finished.
33     print ending time-stamp to log file.
34
35     exit.
36 }
37
38 }
39
40 ▼ run the (task)(number of iterations){
41
42     print to main log file that (task) has started
43     print to individual log file that (task) has started
44
45 ▼     (loop from 0 - number of iterations){
46
47         plaintext = Concatenate "HASHSTRING"+IterationNumber";
48         Hashtext = Hash(plaintext);
49         write Hashtext to individual logfile.
50
51     }
52
53     print to individual log file that (task) has finished.
54     print to main log file that (task) has finished.
55 }
```

Threads:

```
1 PsuedoCode: Tasks
2
3 //threadObject structure
4 threadObject= {Name of the thread: "Thread1", Number of iterations: number of times to hash}
5
6 Main(number of processes, number of iterations){
7
8     - Write to main log file the starting time.
9
10    - Create the thread objects that contain name of thread, and number of iterations.
11
12    - Create the threads, arguments to do runTask and supply the thread object.
13
14    - join each thread to make sure that the program waits till they finish.
15
16    - Write to main log file timestamp once all threads completed.
17
18 }
19
20 run the (task)(threadObject){
21
22    - do a sched yield to make sure that we yield to other processes once this one finishes.
23
24    - lock resources to execute thread.
25
26
27    print to main log file that (task) has started
28    print to individual log file that (task) has started
29
30    (loop from 0 - number of iterations){
31
32        plaintext = Concatenate "HASHSTRING"+IterationNumber";
33        Hashtext = Hash(plaintext);
34        write Hashtext to individual logfile.
35
36    }
37
38    print to individual log file that (task) has finished.
39    print to main log file that (task) has finished.
40
41    - unlock resources
42
43 }
44
45 |
```

**Program Requirements:**

1. openssl-devel-1:1.0.1j-1.fc21.x86\_64

**Execution Instructions**

To execute the experiment as it was done for gathering data.

1. Install required libraries
  - a. *Yum install openssl-devel*
  - b. *Yum install openssl-libs-1.0.1j-1.fc21.x86\_64 zlib-1.2.8-7.fc21.x86\_64*
2. Compile source code.
  - a. *CD ./Processes*
  - b. *Make ProcessRunner*
  - c. *CD ../Threads*
  - d. *Make ThreadRunner*
3. Within main directory, execute
  - a. *SH ExecuteExperiment.sh 4*

To only run the Thread or ProcessRunner

1. Install required libraries
  - a. *Yum install openssl-devel*
  - b. *Yum install openssl-libs-1.0.1j-1.fc21.x86\_64 zlib-1.2.8-7.fc21.x86\_64*
2. Compile with following flags
  - a. *-lcrypto -pthread*

**Testing**

#	Name	Resources	Expected	Actual	Fig
1	Hash specified amount of times	ProcessRunner.c, ThreadRunner.c	The end of the log file to contain the last hash number which should be the same as the iteration number	Same as expected	1.1
2	Generate unique hashes.	ThreadOutputFile X.txt, ProcessOutputFile X.txt	The hashes should all be different from each other.	As expected	4.1
3	Spawn 5 Threads/Processes	ThreadRunner.c, ProcessRunner.c	1 parent thread and 5 child threads.	Same as expected.	2.1A, 2.1B
4	Wait for child processes / other threads.	ProcessRunner.c, ThreadRunner.c	Program will not finish (print out end time) until tasks and processes have all finished writing to file.	Same as expected.	3.1A, 3.1B

```

Test 3 THREADS Case 4 :10000000 hash iterations FINISHED

tail: cannot open './ThreadFiles/ThreadTaskOutputFileThread0.txt' for
ITERATION 10000000 | END TIME: Sun Jan 25 00:23:52 2015

Thread1 END: TIME: Sun Jan 25 00:23:52 2015
ITERATION 10000000 | END TIME: Sun Jan 25 00:23:45 2015

Thread2 END: TIME: Sun Jan 25 00:23:45 2015
ITERATION 10000000 | END TIME: Sun Jan 25 00:24:30 2015

Thread3 END: TIME: Sun Jan 25 00:24:30 2015
ITERATION 10000000 | END TIME: Sun Jan 25 00:24:28 2015

Thread4 END: TIME: Sun Jan 25 00:24:28 2015

```

**Fig 1.1**

**Hash specified amount of times.** ("Iteration 10000000 | END TIME ...") is the output of a "tail -n4" command reading from a log file. This proves that 1000000 iterations were written.

KiB Swap: 839676 total, 387708 free, 451968 used. 245408 avail Mem

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
50536	root	20	0	14684	300	0	D	66.1	0.0	0:06.01	ProcessRunner
50532	root	20	0	14684	300	0	D	65.8	0.0	0:06.01	ProcessRunner
50535	root	20	0	14684	300	0	R	65.8	0.0	0:05.90	ProcessRunner
50533	root	20	0	14684	300	0	D	58.1	0.0	0:05.75	ProcessRunner
50534	root	20	0	14684	300	0	D	56.5	0.0	0:05.97	ProcessRunner
187	root	20	0	0	0	0	S	0.0	0.0	0:01.00	ksuend0

Fix 2.1A 5 child processes running

root	50290	2	50290	0	1	00:16	?			00:00:00	[kworker/0:0]
root	50372	50064	50372	0	6	00:17	pts/3			00:00:00	./ThreadRunner 10000000
root	50372	50064	50374	24	6	00:17	pts/3			00:00:49	./ThreadRunner 10000000
root	50372	50064	50375	25	6	00:17	pts/3			00:00:50	./ThreadRunner 10000000
root	50372	50064	50376	22	6	00:17	pts/3			00:00:46	./ThreadRunner 10000000
root	50372	50064	50377	23	6	00:17	pts/3			00:00:46	./ThreadRunner 10000000
root	50372	50064	50378	23	6	00:17	pts/3			00:00:47	./ThreadRunner 10000000
root	50379	48644	50379	0	1	00:20	pts/2			00:00:00	ps -eLf

Fig 2.1B 5 child threads running.

```

Test 3 THREADS Case 4 :10000000 hash iterations
rm: cannot remove './ThreadFiles/*.txt': No such file or directory
** THREADS PROGRAM START** TIME: Sun Jan 25 00:17:26 2015

** Thread1 START** TIME Sun Jan 25 00:17:26 2015

** Thread5 START** TIME Sun Jan 25 00:17:26 2015

** Thread2 START** TIME Sun Jan 25 00:17:26 2015

** Thread3 START** TIME Sun Jan 25 00:17:26 2015

** Thread4 START** TIME Sun Jan 25 00:17:26 2015

** Thread2 END** TIME Sun Jan 25 00:23:45 2015

** Thread1 END** TIME Sun Jan 25 00:23:52 2015

** Thread5 END** TIME Sun Jan 25 00:24:27 2015

** Thread4 END** TIME Sun Jan 25 00:24:28 2015

** Thread3 END** TIME Sun Jan 25 00:24:30 2015

Threads Complete
** THREADS PROGRAM END** TIME: Sun Jan 25 00:24:30 2015

```

Fig 3.1A Waiting for child threads to finish before exiting.

```

Test 1 PROCESSES Case 3 :1000000 hash iterations
rm: cannot remove './ProcessFiles/*.txt': No such file or directory
** PROCESSES PROGRAM WITH NUMBER OF ITERATIONS: 1000000 START TIME: Sun Jan 25 00:35:31 2015

** PROCESS: 50516 START** TIME Sun Jan 25 00:35:31 2015
** PROCESS: 50517 START** TIME Sun Jan 25 00:35:31 2015

** PROCESS: 50518 START** TIME Sun Jan 25 00:35:31 2015

** PROCESS: 50519 START** TIME Sun Jan 25 00:35:31 2015

** PROCESS: 50520 START** TIME Sun Jan 25 00:35:31 2015

** PROCESS: 50517 END** TIME Sun Jan 25 00:35:43 2015

** PROCESS: 50520 END** TIME Sun Jan 25 00:35:44 2015

** PROCESS: 50519 END** TIME Sun Jan 25 00:35:44 2015

** PROCESS: 50518 END** TIME Sun Jan 25 00:35:44 2015

** PROCESS: 50516 END** TIME Sun Jan 25 00:35:44 2015

** PROGRAM END TIME, Sun Jan 25 00:35:44 2015

```

**Fig 3.1B Waiting for child processes to finish before exiting.**

```

PLAINTEXT: HASHSTRING5789245
HASH:
fffffd2337810fffff6fffffdfffffa833fffffa6fffffe9fffffc41d56fffff8efffff9e31fffff99fffffcf
ffffa23a
ITERATION 5789245 | END TIME: Sun Jan 25 01:20:28 2015

ITERATION 5789246 | TIME: Sun Jan 25 01:20:28 2015

PLAINTEXT: HASHSTRING5789246
HASH:
37fffffa68fffff4fffffaa9684577fffffc04dfffffeaffffffdfffff3834cfffffaa9ffffb6
ITERATION 5789246 | END TIME: Sun Jan 25 01:20:28 2015

ITERATION 5789247 | TIME: Sun Jan 25 01:20:28 2015

PLAINTEXT: HASHSTRING5789247
HASH:
e4fffff8ac766affffff343fffffb73246fffffd4fffffd0227effffffe76fffffe3e5d
ITERATION 5789247 | END TIME: Sun Jan 25 01:20:28 2015

ITERATION 5789248 | TIME: Sun Jan 25 01:20:28 2015

PLAINTEXT: HASHSTRING5789248
HASH:
fffffbf4352fffffb5473fffffb04bfffffe64a7bfffffa556468fffffd45fffffaeffffffe8
ITERATION 5789248 | END TIME: Sun Jan 25 01:20:28 2015

ITERATION 5789249 | TIME: Sun Jan 25 01:20:28 2015

PLAINTEXT: HASHSTRING5789249
HASH:
355bfffffe6fffff8cfffffa6465c1fffff2fffffe0fffffa22e28fffffc2fffff040fffffdcfffffc1fffffb8
1c
ITERATION 5789249 | END TIME: Sun Jan 25 01:20:28 2015

ITERATION 5789250 | TIME: Sun Jan 25 01:20:28 2015

PLAINTEXT: HASHSTRING5789250
HASH:
fffffbf8f38fffff9effffffd8fffffc2fffffb7fffffdfffffb56f6effffffe9763959fffffa3dfffffc3161dfffff
90
ITERATION 5789250 | END TIME: Sun Jan 25 01:20:28 2015

ITERATION 5789251 | TIME: Sun Jan 25 01:20:28 2015

PLAINTEXT: HASHSTRING5789251
HASH:
fffff9e31fffff0fffffa932fffffd4fffffd61f231bfffffb9d35fffffaa9ffffba561fffffb9effffffb6
fffffd5
ITERATION 5789251 | END TIME: Sun Jan 25 01:20:28 2015

```

**Fig 4.1 Generate unique hashes**



**Annex A: Experiment Data**Environment 1 (BCIT Lab) Data

	Iterations				
Processes	1000	10000	100000	1000000	10000000
Average Time	0:00:00	0:00:00	0:00:02	0:00:13	0:02:12
Test 1	0:00:00	0:00:01	0:00:01	0:00:14	0:02:11
Test 2	0:00:00	0:00:00	0:00:02	0:00:13	0:02:10
Test 3	0:00:00	0:00:00	0:00:02	0:00:13	0:02:14

	Iterations				
Threads	1000	10000	100000	1000000	10000000
Average Time	0:00:00	0:00:00	0:00:04	0:00:37	0:06:21
Test 1	0:00:00	0:00:00	0:00:04	0:00:38	0:06:26
Test 2	0:00:00	0:00:00	0:00:04	0:00:37	0:06:21
Test 3	0:00:00	0:00:00	0:00:04	0:00:37	0:06:17

Graph Data

Graph Data	Iterations				
Average time over 3 tests	1000	10000	100000	1000000	10000000
Processes	0:00:00	0:00:00	0:00:02	0:00:13	0:02:12
Threads	0:00:00	0:00:00	0:00:04	0:00:37	0:06:21

**Fig 1B: Data of performance times over amount of iterations hashed in Environment 1 (BCIT Lab)**

Environment 1 (Linux VM) Data

	Iterations				
Processes	1000	10000	100000	1000000	10000000
Average Time	0:00:00	0:00:00	0:00:02	0:00:13	0:02:18
Test 1	0:00:00	0:00:00	0:00:02	0:00:13	0:02:20
Test 2	0:00:00	0:00:00	0:00:01	0:00:13	0:02:17
Test 3	0:00:00	0:00:00	0:00:02	0:00:13	0:02:18

	Iterations				
Threads	1000	10000	100000	1000000	10000000
Average Time	0:00:00	0:00:01	0:00:04	0:00:41	0:06:42
Test 1	0:00:01	0:00:00	0:00:04	0:00:41	0:06:40
Test 2	0:00:00	0:00:01	0:00:04	0:00:44	0:06:57
Test 3	0:00:00	0:00:01	0:00:03	0:00:39	0:06:30

Graph Data

Graph Data	Iterations				
Average time over 3 tests	1000	10000	100000	1000000	10000000
Processes	0:00:00	0:00:00	0:00:02	0:00:13	0:02:18
Threads	0:00:00	0:00:01	0:00:04	0:00:41	0:06:42

**Fig 1D: Data of performance times over amount of iterations hashed in Environment 2 (Linux VM)**